OPTIMIZED FOR

Hyper-Converged

Power and flexibility are the key to any modern infrastructure setup. By bringing together the best of flexible open hardware, cost-effective scaling and the efficient, easy-to-use model, Sesame optimized for hyper-converged infrastructure is perfect for enterprises looking for an alternative to traditional silo approaches.

Sesame provides a converged infrastructure layer making it easy to scale compute, storage and networking together using proven hyperscale technology and common building blocks.

With our stable and scalable infrastructure, deploying apps is easier than ever. Configurable clusters that move and scale as your needs change make updating code, applications and data structures straightforward for new or for legacy applications.

Key Benefits

- **Simplicity** of standardized, converged nodes

- **Optimized for converged** infrastructure deployments

- **Compute, storage and network capability** can scale in lock-step as needed

- **Fast storage** with the ability to store and process data at scale, from modest to large data sets

- **Fully open** configuration and manageability tools, qualified and optimized for Linux

Available three- or six-year warranty

From 6 to 30 converged nodes

Resilient performance via proven rack-level architecture

Internal Switches
Ingress/Egress

Converged Nodes
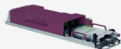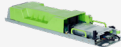
Infrastructure/
Management Nodes

Example shown: **Scale**

Configuration

	BASE **12 Nodes**	GROWTH **18 Nodes**	MEDIUM **24 Nodes**	SCALE **36 Nodes**
Converged Nodes	6	12	18	30
Infrastructure / Management Nodes	3	3	3	3
Ingress / Egress Nodes	3	3	3	3
External Switches	Redundant 25G (50/100G uplink)			

Performance Specs

- From six to 30 converged compute and storage nodes

- Dual-socket compute and terabytes of flash storage with 25 GbE connectivity

- Configurable, cost-effective management and infrastructure nodes in groups of three

- Redundant top-of-rack switches and uplinks compatible with your existing network environments

Node Specs

	CONFIGURATION	CATEGORY
Compute 2S	Optimized for two-socket converged compute + storage	High CPU + fast storage
Management, Ingress / Egress, Infrastructure	Optimized for infrastructure services	Modest CPU + modest storage

Check out these other Sesame solutions built and tuned for what you need your infrastructure to do

Optimized for Kubernetes
Optimized for open source, cloud-native workloads requiring containerization and Kubernetes orchestration.

Optimized for AI/ML
Workloads that require massive data volumes in demanding learning environments using models like TensorFlow.

Optimized for Edge Computing
Optimized for deployment in distributed environments that require smaller form factors running forward-deployed workloads at scale.